

Contents

Storage Architecture - Tri-Store Hybrid System	1
Overview	1
Storage Decision Matrix	2
Data Flow Diagrams	2
1. CSV Upload Flow	2
2. Supporting Document Flow	2
3. Validation Flow	3
4. Analytics Query Flow	3
5. RAG Chatbot Flow	3
PostgreSQL Schema	4
Tables	4
Indexes	5
MongoDB Collections	5
supporting_docs	5
audit_trail	6
validation_results	6
File System Structure	7
Technology Stack	7
Database Drivers	7
Storage Libraries	8
Connection Management	8
PostgreSQL	8
MongoDB	8
File Storage	8
Environment Variables	8
Docker Setup	9
Initialization	9
Performance Optimizations	9
Backup & Recovery	9
PostgreSQL	9
MongoDB	10
Files	10
Phase Alignment	10

Storage Architecture - Tri-Store Hybrid System

Overview

Project Aura uses a **tri-store architecture** that leverages the strengths of three different storage systems:

Architecture Overview

1. **PostgreSQL** - Structured financial data

2. **MongoDB** - Semi-structured metadata and audit logs
3. **File System** - Unstructured data and binary files

This hybrid approach optimizes for performance, flexibility, and developer experience.

Storage Decision Matrix

Data Type	PostgreSQL	MongoDB	Files	Rationale
GL Account Balances	Primary		Parquet cache	Structured, requires joins, SQL
User Information	Primary			Relational, foreign keys
Responsibility Matrix	Primary			Many-to-many relationships
Review Status	Primary			Transactional updates
Supporting Doc Metadata		Primary		Flexible schema, varies per document
Actual PDFs/Excel			Primary	Binary data, large files
Comments & Annotations		Primary		Nested structure, replies
Audit Trail Events	Summary	Detailed		Both: Summary in PG, full log in MongoDB
Validation Results	Pass/Fail	Full Report	JSON backup	PG for status, Mongo for detailed results
Vector Embeddings			ChromaDB	Specialized vector storage
ML Model Artifacts			MLflow	Binary model files
Raw CSV Uploads			Landing zone	Temporary storage before processing
Processed DataFrames			Parquet	Fast columnar access

Data Flow Diagrams

1. CSV Upload Flow

```

User uploads CSV (Streamlit)
  ↓
Save to data/raw/ (File System)
  ↓
Parse & Validate Schema
  ↓
Load to PostgreSQL (gl_accounts table)
  ↓
Log metadata to MongoDB (upload_tracking)
  ↓
Cache as Parquet in data/processed/ (File System)

```

2. Supporting Document Flow

User uploads PDF/Excel

```
↓  
Save binary to data/supporting_docs/{gl_code}/ (File System)  
↓  
Save metadata to MongoDB (supporting_docs collection)  
{  
    gl_code: "1001",  
    files: [{name, path, uploaded_by, timestamp}],  
    comments: []  
}  
↓  
Log audit event to MongoDB (audit_trail)  
↓  
Update review_log in PostgreSQL
```

3. Validation Flow

```
Trigger validation (Button click)  
↓  
Load data from PostgreSQL  
↓  
Run Great Expectations suite  
↓  
Save pass/fail to PostgreSQL (review_log)  
↓  
Save full results to MongoDB (validation_results)  
↓  
Save checkpoint to data/processed/ (File System)
```

4. Analytics Query Flow

```
User requests dashboard  
↓  
Check Parquet cache in data/processed/  
    HIT: Load from Parquet  
    MISS: Query PostgreSQL  
        ↓  
        Compute aggregations (SQL)  
        ↓  
        Cache results as Parquet  
    ↓  
Generate Plotly charts  
↓  
Display in Streamlit
```

5. RAG Chatbot Flow

```
User asks question
```

```

    ↓
Embed query using OpenAI
    ↓
Vector search in ChromaDB (data/vectors/)
    ↓
Retrieve relevant context
    ↓
Agent decides: Call tool OR use context
        Tool call: Query PostgreSQL for live data
        Context: Answer from FAQ docs
    ↓
Generate response with LangChain
    ↓
Log interaction to MongoDB (audit_trail)

```

PostgreSQL Schema

Tables

`users`

```

id SERIAL PRIMARY KEY
name VARCHAR(255) NOT NULL
email VARCHAR(255) UNIQUE NOT NULL
department VARCHAR(100)
role VARCHAR(50)
created_at TIMESTAMP DEFAULT NOW()
updated_at TIMESTAMP DEFAULT NOW()

```

`gl_accounts`

```

id SERIAL PRIMARY KEY
account_code VARCHAR(50) NOT NULL
account_name VARCHAR(255) NOT NULL
entity VARCHAR(255) NOT NULL
balance DECIMAL(18, 2) NOT NULL
period VARCHAR(20) NOT NULL
assigned_user_id INTEGER REFERENCES users(id)
review_status VARCHAR(50) DEFAULT 'pending'
created_at TIMESTAMP DEFAULT NOW()
updated_at TIMESTAMP DEFAULT NOW()
UNIQUE(account_code, entity, period)

```

`responsibility_matrix`

```

id SERIAL PRIMARY KEY
gl_code VARCHAR(50) NOT NULL
user_id INTEGER REFERENCES users(id)
role VARCHAR(50) NOT NULL
created_at TIMESTAMP DEFAULT NOW()
UNIQUE(gl_code, user_id, role)

review_log

id SERIAL PRIMARY KEY
gl_account_id INTEGER REFERENCES gl_accounts(id)
reviewer_id INTEGER REFERENCES users(id)
status VARCHAR(50) NOT NULL
comments TEXT
reviewed_at TIMESTAMP DEFAULT NOW()

```

Indexes

- idx_gl_accounts_period on gl_accounts(period)
 - idx_gl_accounts_entity on gl_accounts(entity)
 - idx_gl_accounts_status on gl_accounts(review_status)
 - idx_review_log_gl on review_log(gl_account_id)
-

MongoDB Collections

```

supporting_docs

{
  _id: ObjectId,
  gl_code: "1001",
  period: "2025-01",
  entity: "Adani Ports",
  files: [
    {
      name: "invoice_jan.pdf",
      path: "data/supporting_docs/1001/invoice_jan.pdf",
      uploaded_by: "john@adani.com",
      uploaded_at: ISODate("2025-01-15")
    }
  ],
  comments: [
    {
      user: "Manager",
      text: "Please clarify this variance",
      timestamp: ISODate("2025-01-16"),
    }
  ]
}

```

```
        replies: []
    }
],
created_at: ISODate,
updated_at: ISODate
}
```

Indexes:

- gl_code
- (gl_code, period) compound

audit_trail

```
{
  _id: ObjectId,
  gl_code: "1001",
  action: "reviewed",
  actor: {
    user_id: "U123",
    name: "Jane Smith",
    email: "jane@adani.com"
  },
  details: {
    previous_status: "pending",
    new_status: "approved"
  },
  timestamp: ISODate
}
```

Indexes:

- gl_code
- timestamp (descending)
- (gl_code, timestamp) compound

validation_results

```
{
  _id: ObjectId,
  gl_code: "1001",
  period: "2025-01",
  validation_suite: "trial_balance_suite",
  results: [
    expectations: [...],
    success_count: 45,
    failure_count: 2,
    failures: [...]
  ]
}
```

```
    },
    passed: false,
    validated_at: ISODate
}
```

Indexes:

- gl_code
 - validation_suite
-

File System Structure

```
data/
  raw/                                # Landing zone for CSV uploads
    trial_balance_jan.csv
    trial_balance_feb.csv

  processed/                            # Cached Parquet files
    gl_accounts_2025-01.parquet
    analytics_summary.parquet

  supporting_docs/                     # Organized by GL code
    1001/
      invoice_jan.pdf
      reconciliation.xlsx
    2005/
      backup_working.pdf

  vectors/                             # ChromaDB persistence
    chroma.sqlite3

  sample/                             # Sample data for testing
    trial_balance_sample.csv
```

Technology Stack

Database Drivers

- **psycopg2-binary** 2.9.9 - PostgreSQL adapter
- **SQLAlchemy** 2.0.30 - ORM and query builder
- **pymongo** 4.6.2 - MongoDB driver
- **pyarrow** 15.0.0 - Parquet file support

Storage Libraries

- pandas - DataFrame operations
 - chromadb - Vector storage for RAG
 - mlflow - ML model artifact storage
-

Connection Management

PostgreSQL

```
from src.db import get_postgres_session

session = get_postgres_session()
users = session.query(User).all()
session.close()
```

MongoDB

```
from src.db import get_mongo_database

db = get_mongo_database()
collection = db["supporting_docs"]
doc = collection.find_one({"gl_code": "1001"})
```

File Storage

```
from src.db.storage import save_raw_csv, load_processed_parquet

# Save CSV
save_raw_csv(df, "trial_balance.csv")

# Load Parquet cache
cached_df = load_processed_parquet("analytics_summary")
```

Environment Variables

```
# PostgreSQL
POSTGRES_HOST=localhost
POSTGRES_PORT=5432
POSTGRES_DB=finnovate
POSTGRES_USER=admin
POSTGRES_PASSWORD=hackathon2025

# MongoDB
```

```
MONGO_HOST=localhost  
MONGO_PORT=27017  
MONGO_DB=finnovate
```

Docker Setup

Start databases:

```
docker-compose up -d
```

Check status:

```
docker-compose ps
```

View logs:

```
docker-compose logs -f postgres  
docker-compose logs -f mongodb
```

Initialization

Run bootstrap script to:

1. Start Docker containers
2. Create PostgreSQL tables
3. Initialize MongoDB collections with indexes
4. Create file system directories

```
.\scripts\bootstrap.ps1
```

Performance Optimizations

1. **PostgreSQL Indexes** - Fast queries on period, entity, status
 2. **Parquet Caching** - Columnar format for analytics
 3. **MongoDB Compound Indexes** - Fast lookups on (gl_code, period)
 4. **Connection Pooling** - Reuse database connections
 5. **ChromaDB Persistence** - Avoid re-embedding on startup
-

Backup & Recovery

PostgreSQL

```
docker exec finnovate-postgres pg_dump -U admin finnovate > backup.sql
```

MongoDB

```
docker exec finnovate-mongodb mongodump --out /backup
```

Files

Regular filesystem backups of `data/` directory.

Phase Alignment

- **Phase 1 (Day 1):** File storage + basic PostgreSQL (`gl_accounts`, `users`)
- **Phase 2 (Day 2):** Full PostgreSQL schema + MongoDB (`validation_results`)
- **Phase 3 (Day 3):** MongoDB `audit_trail` + `supporting_docs`
- **Phase 4 (Day 4):** ChromaDB vector store for RAG
- **Phase 5 (Day 5-6):** Performance tuning, caching, observability